



ARMv4 Memory Instructions

Instructor: Dr. Vinicius Prado da Fonseca (vpradodafons@online.mun.ca)

Memory

- A computer's memory is controlled by address lines, data lines, a read/write control line, and a clock line (Lec 17)
- The address lines selects the memory location to access
- The data lines can be used to send or receive data to/from the memory
- The read/write control line determine the direction data is being sent, and the clock line controls when the memory operation occurs

Memory

- Memory can be treated like a byte array
- The address lines act as an index into the array
- Thus, memory reads and writes can be described by:

```
r0 = mem[add]; // read to r0 from memory using add as the index
```

```
mem[add] = r0; // write to memory from r0 using add as the index
```

ARMv4 Memory Load Instructions

- Memory load instructions transfer the contents of a memory location to a CPU register
- A byte or a 32-bit word can be transferred
- Load instructions have the following format:

```
ldr Rd, [Rn, #immed]           // R[d] <= M[ R[n] + #immed]
```

```
ldr Rd, [Rn, Rm, sh #immed]    // R[d] <= M[ R[n] + (R[m] sh #immed)]
```

```
ldrb Rd, [Rn, #immed]          // R[d] <= M[ R[n] + #immed]
```

```
ldrb Rd, [Rn, Rm, sh #immed]   // R[d] <= M[ R[n] + (R[m] sh #immed)]
```

Memory Load Instructions

- ldr loads 32 bits from memory into a 32 bit register
- ldrb loads 8 bits from memory into a 32 bit register
 - The upper 24 bits are set to zero
- Similar to the data processing instruction has second operand:
 - an immediate value
 - a register with an optional shift/rotate
- In the immediate offset, the address is calculated by adding Rn with the immediate value
- In the register offset, Rn and Rm are added

Memory Load Instructions

- Rm can also be
 - left shifted
 - right shifted
 - arithmetic right shifted
 - right rotated before the addition
- The Rd register will contain the data, and the Rn and Rm registers are used to calculate the address
- The address is form by either $Rn + Rm$, or $Rn - Rm$
- An immediate value can be added or subtracted to Rn to form the address

Memory Store Instructions

- str stores 32 bits to memory from a register
- strb stores 8 bits to memory from the lower 8 bits of a register
- Similar to the data processing and load instructions, an immediate value or a register with an optional shift/rotate is available as the second operand

```
str Rd, [Rn, #immed]           // M[ R[n] + #immed] <= R[d]
```

```
str Rd, [Rn, Rm, sh #immed]    // M[ R[n] + (R[m] sh #immed) <= R[d]
```

```
strb Rd, [Rn, #immed]          // M[ R[n] + #immed] <= R[d]
```

```
strb Rd, [Rn, Rm, sh #immed]   // M[ R[n] + (R[m] sh #immed) <= R[d]
```

Memory Address Calculations

- The way the memory address is calculated is often called the addressing mode
- The load and store instructions provide four different ways of calculating the memory address
- The simplest is:

```
ldr r2, [r3]
```

```
str r2, [r3]
```

- The value of r3 contains the address
- If r3 is 0x2000, then the location accessed is 0x2000
 - $r2 \leftarrow M[0x2000]$

Memory Address Calculations

- A fixed constant can be added to the register to give the location
- The constant is between -4095 and 4095 for byte and 32-bit memory accesses

```
ldr r2, [r3, #20]  
str r2, [r3, #20]
```

- The location accessed is calculated by adding the contents of r3 with the constant 20 or 0x14 in hexadecimal
- If r3 is 0x400, then the address is 0x414
 - $r2 \leq M[0x414]$
 - $M[0x414] \leq r2$

Memory Address Calculations

- This method and the first are identical
- The first method uses a constant 0
- Thus the first method could have been written as:

```
ldr r2, [r3,#0]
```

```
str r2, [r3,#0]
```

Memory Address Calculations

- Two registers can be used in the address calculation

```
ldr r2, [r3,r1]
```

```
str r2, [r3,r1]
```

- The contents of r3 and r1 are added to give the location to access
- If r3 is 0x2000 and r1 is 0x100, then the location is 0x2100

Memory Address Calculations

- Subtraction can be specified with:

```
ldr r2, [r3,-r1]
```

```
str r2, [r3,-r1]
```

- The r1 register is subtracted from r3 to yield the address
- If r3 is 0x2000 and r1 is 0x100, then the location is 0x1f00

Memory Address Calculations

- In the fourth method, the second register is shifted left, shifted right, or rotated by a fixed amount and the added to the first register
- Some examples are:

```
ldr r2, [r3, r1, lsl #2]
```

```
str r2, [r3, r1, lsl #2]
```

Memory Address Calculations

```
ldr r2, [r3, r1, lsl #2]
```

```
str r2, [r3, r1, lsl #2]
```

- Assuming
 - r3 contains 0x800
 - r1 is 0x8
- Location accessed is calculated by shifting “r1” 2 position to the left (i.e., multiplying by 4)
- Then adding r3
- The value is $0x800 + 4 * 0x8$, which is 0x820
- $r2 \leq M[0x820]$ or $M[0x820] \leq r2$
- This last method is often used to implement array access.

Sum Array With 32-bit Integers (sum_array.S)

- r0 = address to the start of the array
- r1 = number of elements to sum
- Moves r0 to r2
 - r2 now contains the start of the array
- r0 is used to hold the sum
- r0 is set to zero with eor r0, r0, r0
- r3 holds the value loaded from memory
- The loop sums the array entries from the end index to the start

```
sum_array:
    mov    r2, r0
    eor    r0, r0, r0    // r0 = 0
top:
    subs   r1, r1, #1
    blt    done
    ldr    r3, [r2, r1, lsl #2]
    add    r0, r0, r3
    b      top
done:
    mov    pc, lr
```

Sum Array With 32-bit Integers (sum_array.S)

- `subs r1, r1, #1` subtracts 1 from the count
- If the count is **less than 0**
 - loop is terminated by `blt done`
- An entry in the array is loaded with:

`ldr r3, [r2, r1, lsl #2]`

```
sum_array:
    mov    r2, r0
    eor    r0, r0, r0    // r0 = 0

top:
    subs   r1, r1, #1
    blt    done
    ldr    r3, [r2, r1, lsl #2]
    add    r0, r0, r3
    b      top

done:
    mov    pc, lr
```


Sum Array With 32-bit Integers (sum_array.S)

- The index, r1 is scaled by 4 to account for the 32 bits
 - 4 bytes per each integer in the array
 - lsl #2 scales by 4
- The address is calculated by $r2 + r1 * 4$
 - r2 is base address
 - $r1 * 4$ is next number address
- r3 is added to r0 to form the sum
- The code then branches to the top of the loop to load the next entry
- Test with “sum_array_main.c”

```
arm-linux-gnueabi-gcc --static  
sum_array_main.c sum_array.S
```

```
sum_array:  
    mov    r2, r0  
    eor    r0, r0, r0    // r0 = 0  
  
top:  
    subs   r1, r1, #1  
    blt    done  
    ldr     r3, [r2, r1, lsl #2]  
    add    r0, r0, r3  
    b      top  
  
done:  
    mov    pc, lr
```

Sum Array With 32-bit Integers

- A C-language version of sum_array
- Hand written assembler is very similar to the compiler generated code

```
arm-linux-gnueabi-gcc -march=armv4 -Os -S sum_array.c
```

```
int sum_array( int arr[], int num ) {  
    int sum = 0;  
    while( --num >= 0 ) {  
        sum += arr[num];  
    }  
    return sum;  
}
```

Sum Array With 32-bit Integers

- More straight forward C solution

```
arm-linux-gnueabi-gcc -march=armv4 -Os -S sum_array_v1.c
```

```
int sum_array( int arr[], int num ) {  
    int sum = 0;  
    for( int i = 0; i < num; i++ ) {  
        sum += arr[i];  
    }  
    return sum;  
}
```

Initializing A Sequence (init_seq.S)

- An array is initialized with the sequence 0, 1, 2, ...
- r0 contains start of the array
- r1 number of entries
- `str r2, [r0, r2, lsl #2]`
 - $M[r0 + 4 * r2] \leq r2$
- `movhs pc, lr` **return** from the routine
- Conditional
 - r2 is higher or the same as r1
- The **condition codes** are set by `cmp r2, r1`

```
init_seq:
    mov    r2, #0
top:
    cmp    r2, r1
    movhs  pc, lr
    str    r2, [r0, r2, lsl #2]
    add    r2, r2, #1
    b      top
```

Dot Product (dot_product.S)

$$[1, 2, 3] \cdot [4, 5, 6] = 1 \times 4 + 2 \times 5 + 3 \times 6 = 32$$

- This function accepts three arguments
 - The location of the array a
 - The location of the array b
 - The number of elements
- Since this function uses r4, r5, and r6, these registers are saved with `push {r4,r5,r6}`
- At the end of this function the values are restored with `pop {r4,r5,r6}`

Dot Product (dot_product.S)

- A loop from top to done is used to calculate the dot product
- Each element is read with a ldr instruction
- The mla r0, r5, r6, r0 calculates $r0 = (r5 * r6) + r0$
- mla stands for multiply and accumulate
- r3 is used to index into the array

top:

```
    cmp    r3, r2
    bhs    done
    ldr    r5, [r4, r3, lsl #2]
    ldr    r6, [r1, r3, lsl #2]
    mla    r0, r5, r6, r0
    add    r3, r3, #1
    b      top
```

done:

```
    pop    {r4, r5, r6}
    mov    pc, lr
```

String Capitalize (`capitalize.S`)

- Strings in C are represented by byte arrays
- The end of the string is indicated by a byte with a value of zero
- This zero byte is called a NUL
- Code in `capitalize.S` copies a string to another array and any lower case letters are converted to uppercase

String Capitalize (capitalize.S)

- Bytes are loaded from memory with the ldrb instruction
- A byte is stored with the strb instruction
- The teq instruction tests if loaded bytes is zero
- The loop is terminated if it is zero
- The byte is then checked to see if it is lower case

top:

```
ldrb r3, [r1, r2]
teq r3, #0
beq done
cmp r3, #'a
blt skip
cmp r3, #'z
bgt skip
bic r3, r3, #0x20
```

skip:

```
strb r3, [r0, r2]
add r2, r2, #1
b top
```

done:

```
strb r3, [r0, r2]
mov pc, lr
```


String Capitalize (capitalize.S)

- The bic r3, r3, #0x20 (bit clear) instruction converts a lower case character to upper case
 - A = 0100 0001
 - a = 0110 0001 AND ~(0010 0000) = 0110 0001 AND 1101 1111 = 0100 0001 = A
- The last strb instruction before mov pc, lr places a zero byte in the destination string

top:

```
ldrb r3, [r1, r2]
teq r3, #0
beq done
cmp r3, #'a
blt skip
cmp r3, #'z
bgt skip
bic r3, r3, #0x20
```

skip:

```
strb r3, [r0, r2]
add r2, r2, #1
b top
```

done:

```
strb r3, [r0, r2]
mov pc, lr
```



Unit 5 done

Next: Micro-Architecture of ARM



Questions?